

Data Science Capstone Project- HealthCare

Problem Statement

NIDDK (National Institute of Diabetes and Digestive and Kidney Diseases) research creates knowledge about and treatments for the most chronic, costly, and consequential diseases. The dataset used in this project is originally from NIDDK. The objective is to predict whether or not a patient has diabetes, based on certain diagnostic measurements included in the dataset. Build a model to accurately predict whether the patients in the dataset have diabetes or not.

Dataset Description

The datasets consists of several medical predictor variables and one target variable (Outcome). Predictor variables includes the number of pregnancies the patient has had, their BMI, insulin level, age, and more.

Variables Description

Pregnancies -----> Number of times pregnant

Glucose -----> Plasma glucose concentration in an oral glucose tolerance test

BloodPressure -----> Diastolic blood pressure (mm Hg)

SkinThickness -----> Triceps skinfold thickness (mm)

Insulin -----> Two hour serum insulin

BMI -----> Body Mass Index

DiabetesPedigreeFunction -----> Diabetes pedigree function

Age -----> Age in years

Outcome -----> Class variable (either 0 or 1). 268 of 768 values are 1, and the others are 0

Project Task: Week 1

Data Exploration:

1. Perform descriptive analysis. Understand the variables and their corresponding values. On the columns below, a value of zero does not make sense and thus indicates missing value:

- Glucose
- BloodPressure
- SkinThickness
- Insulin
- BMI

In [1]:

```
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.neighbors import KNeighborsClassifier
import warnings
warnings.filterwarnings('ignore')
```

In [2]:

```
df= pd.read_csv('health care diabetes.csv')
```

In [3]:

```
df.head()
```

Out[3]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|---|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 0 | 6 | 148 | 72 | 35 | 0 | 33.6 | 0.627 | 50 | 1 |
| 1 | 1 | 85 | 66 | 29 | 0 | 26.6 | 0.351 | 31 | 0 |
| 2 | 8 | 183 | 64 | 0 | 0 | 23.3 | 0.672 | 32 | 1 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 | 0 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 | 1 |

In [4]:

```
df.tail()
```

Out[4]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 | 0 |
| 764 | 2 | 122 | 70 | 27 | 0 | 36.8 | 0.340 | 27 | 0 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 | 0 |

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|---------|
| 766 | 1 | 126 | 60 | 0 | 0 | 30.1 | 0.349 | 47 | 1 |
| 767 | 1 | 93 | 70 | 31 | 0 | 30.4 | 0.315 | 23 | 0 |

In [5]:

```
df.describe()
```

Out[5]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age | Outcome |
|-------|-------------|------------|---------------|---------------|------------|------------|--------------------------|------------|------------|
| count | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 | 768.000000 |
| mean | 3.845052 | 120.894531 | 69.105469 | 20.536458 | 79.799479 | 31.992578 | 0.471876 | 33.240885 | 0.348958 |
| std | 3.369578 | 31.972618 | 19.355807 | 15.952218 | 115.244002 | 7.884160 | 0.331329 | 11.760232 | 0.476951 |
| min | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.000000 | 0.078000 | 21.000000 | 0.000000 |
| 25% | 1.000000 | 99.000000 | 62.000000 | 0.000000 | 0.000000 | 27.300000 | 0.243750 | 24.000000 | 0.000000 |
| 50% | 3.000000 | 117.000000 | 72.000000 | 23.000000 | 30.500000 | 32.000000 | 0.372500 | 29.000000 | 0.000000 |
| 75% | 6.000000 | 140.250000 | 80.000000 | 32.000000 | 127.250000 | 36.600000 | 0.626250 | 41.000000 | 1.000000 |
| max | 17.000000 | 199.000000 | 122.000000 | 99.000000 | 846.000000 | 67.100000 | 2.420000 | 81.000000 | 1.000000 |

In [6]:

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

```
RangeIndex: 768 entries, 0 to 767
```

```
Data columns (total 9 columns):
```

```
#      Column                                Non-Null Count  Dtype
---  -
0      Pregnancies                        768 non-null     int64
1      Glucose                            768 non-null     int64
2      BloodPressure                      768 non-null     int64
3      SkinThickness                     768 non-null     int64
4      Insulin                          768 non-null     int64
5      BMI                              768 non-null     float64
6      DiabetesPedigreeFunction          768 non-null     float64
7      Age                              768 non-null     int64
8      Outcome                          768 non-null     int64
```

```
dtypes: float64(2), int64(7)
```

memory usage: 54.1 KB

In [7]:

```
df.columns
```

Out[7]:

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin',  
      'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],  
      dtype='object')
```

In [8]:

```
df.isnull().sum()
```

Out[8]:

```
Pregnancies      0  
Glucose          0  
BloodPressure    0  
SkinThickness    0  
Insulin          0  
BMI              0  
DiabetesPedigreeFunction  0  
Age              0  
Outcome          0  
dtype: int64
```

In [9]:

```
df.nunique()
```

Out[9]:

```
Pregnancies      17  
Glucose          136  
BloodPressure    47  
SkinThickness    51  
Insulin          186  
BMI              248  
DiabetesPedigreeFunction  517  
Age              52  
Outcome          2  
dtype: int64
```

In [10]:

```
df.shape
```

Out[10]:

```
(768, 9)
```

In [11]:

```
#Checking if there is any 0 value
```

```
df[df[['Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin', 'BMI']]==0].count()
```

Out[11]:

```
Pregnancies      0
```

| | |
|--------------------------|-----|
| Glucose | 5 |
| BloodPressure | 35 |
| SkinThickness | 227 |
| Insulin | 374 |
| BMI | 11 |
| DiabetesPedigreeFunction | 0 |
| Age | 0 |
| Outcome | 0 |

dtype: int64

In [12]:

```
(df[df[['Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI']]==0].count()/len(df))*100
```

Out[12]:

| | |
|--------------------------|-----------|
| Pregnancies | 0.000000 |
| Glucose | 0.651042 |
| BloodPressure | 4.557292 |
| SkinThickness | 29.557292 |
| Insulin | 48.697917 |
| BMI | 1.432292 |
| DiabetesPedigreeFunction | 0.000000 |
| Age | 0.000000 |
| Outcome | 0.000000 |

dtype: float64

In [13]:

```
#Replacing 0 with the median
```

```
for i in ['Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI']:
    print (i, 'Old Median:', df[i].median())
    Median_Value=df[df[i]!=0][i].median()
    print ('New Median',Median_Value, '\n')
    df[i].replace(0,Median_Value,inplace=True)
```

Glucose Old Median: 117.0
New Median 117.0

BloodPressure Old Median: 72.0
New Median 72.0

SkinThickness Old Median: 23.0
New Median 29.0

Insulin Old Median: 30.5
New Median 125.0

BMI Old Median: 32.0
New Median 32.3

In [14]:

```
df[df[['Glucose', 'BloodPressure', 'SkinThickness',
'Insulin', 'BMI']]==0].count()
```

Out[14]:

```
Pregnancies      0
Glucose           0
BloodPressure     0
SkinThickness     0
Insulin           0
BMI               0
DiabetesPedigreeFunction  0
Age               0
Outcome           0
dtype: int64
```


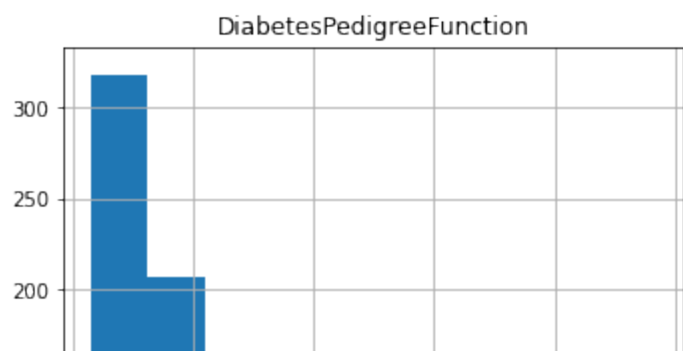

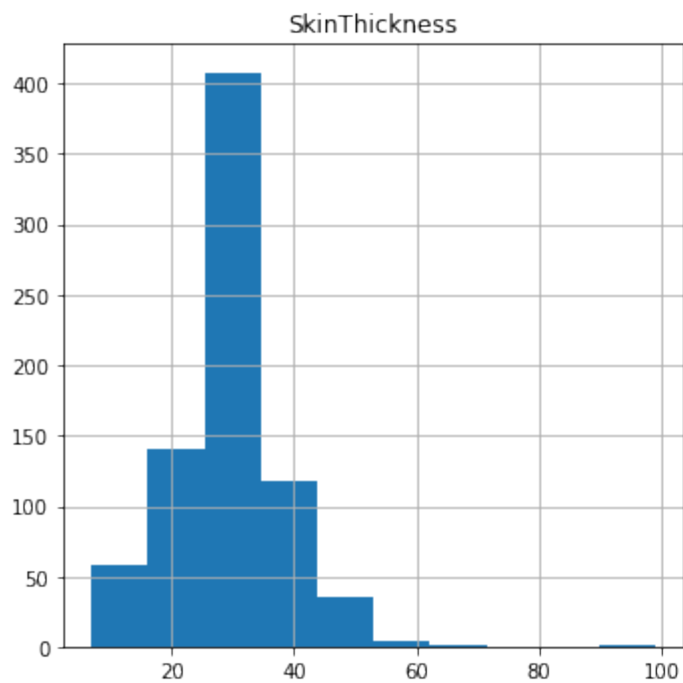


1. 1. Visually explore these variables using histograms. Treat the missing values accordingly.

In [15]:

```
df.hist(figsize=(20,20))
```

Out[15]:

```
array([[<AxesSubplot:title={'center':'Pregnancies'}>,
       <AxesSubplot:title={'center':'Glucose'}>,
       <AxesSubplot:title={'center':'BloodPressure'}>],
       [<AxesSubplot:title={'center':'SkinThickness'}>,
       <AxesSubplot:title={'center':'Insulin'}>,
       <AxesSubplot:title={'center':'BMI'}>],
       [<AxesSubplot:title={'center':'DiabetesPedigreeFunction'}>,
       <AxesSubplot:title={'center':'Age'}>,
       <AxesSubplot:title={'center':'Outcome'}>]], dtype=object)
```




1. There are integer and float data type variables in this dataset. Create a count (frequency) plot describing the data types and the count of variables.

In [16]:

```
df.dtypes.value_counts().plot(kind='bar')
```

Out[16]:

<AxesSubplot:>



Project Task: Week 2

Data Exploration:


1. Check the balance of the data by plotting the count of outcomes by their value. Describe your findings and plan future course of action.

In [17]:

```
df.Outcome.value_counts().plot(kind='bar')
```

Out[17]:

<AxesSubplot:>



The graph shows that the number of patients who are diabetic is half of the patients who are non-diabetic.

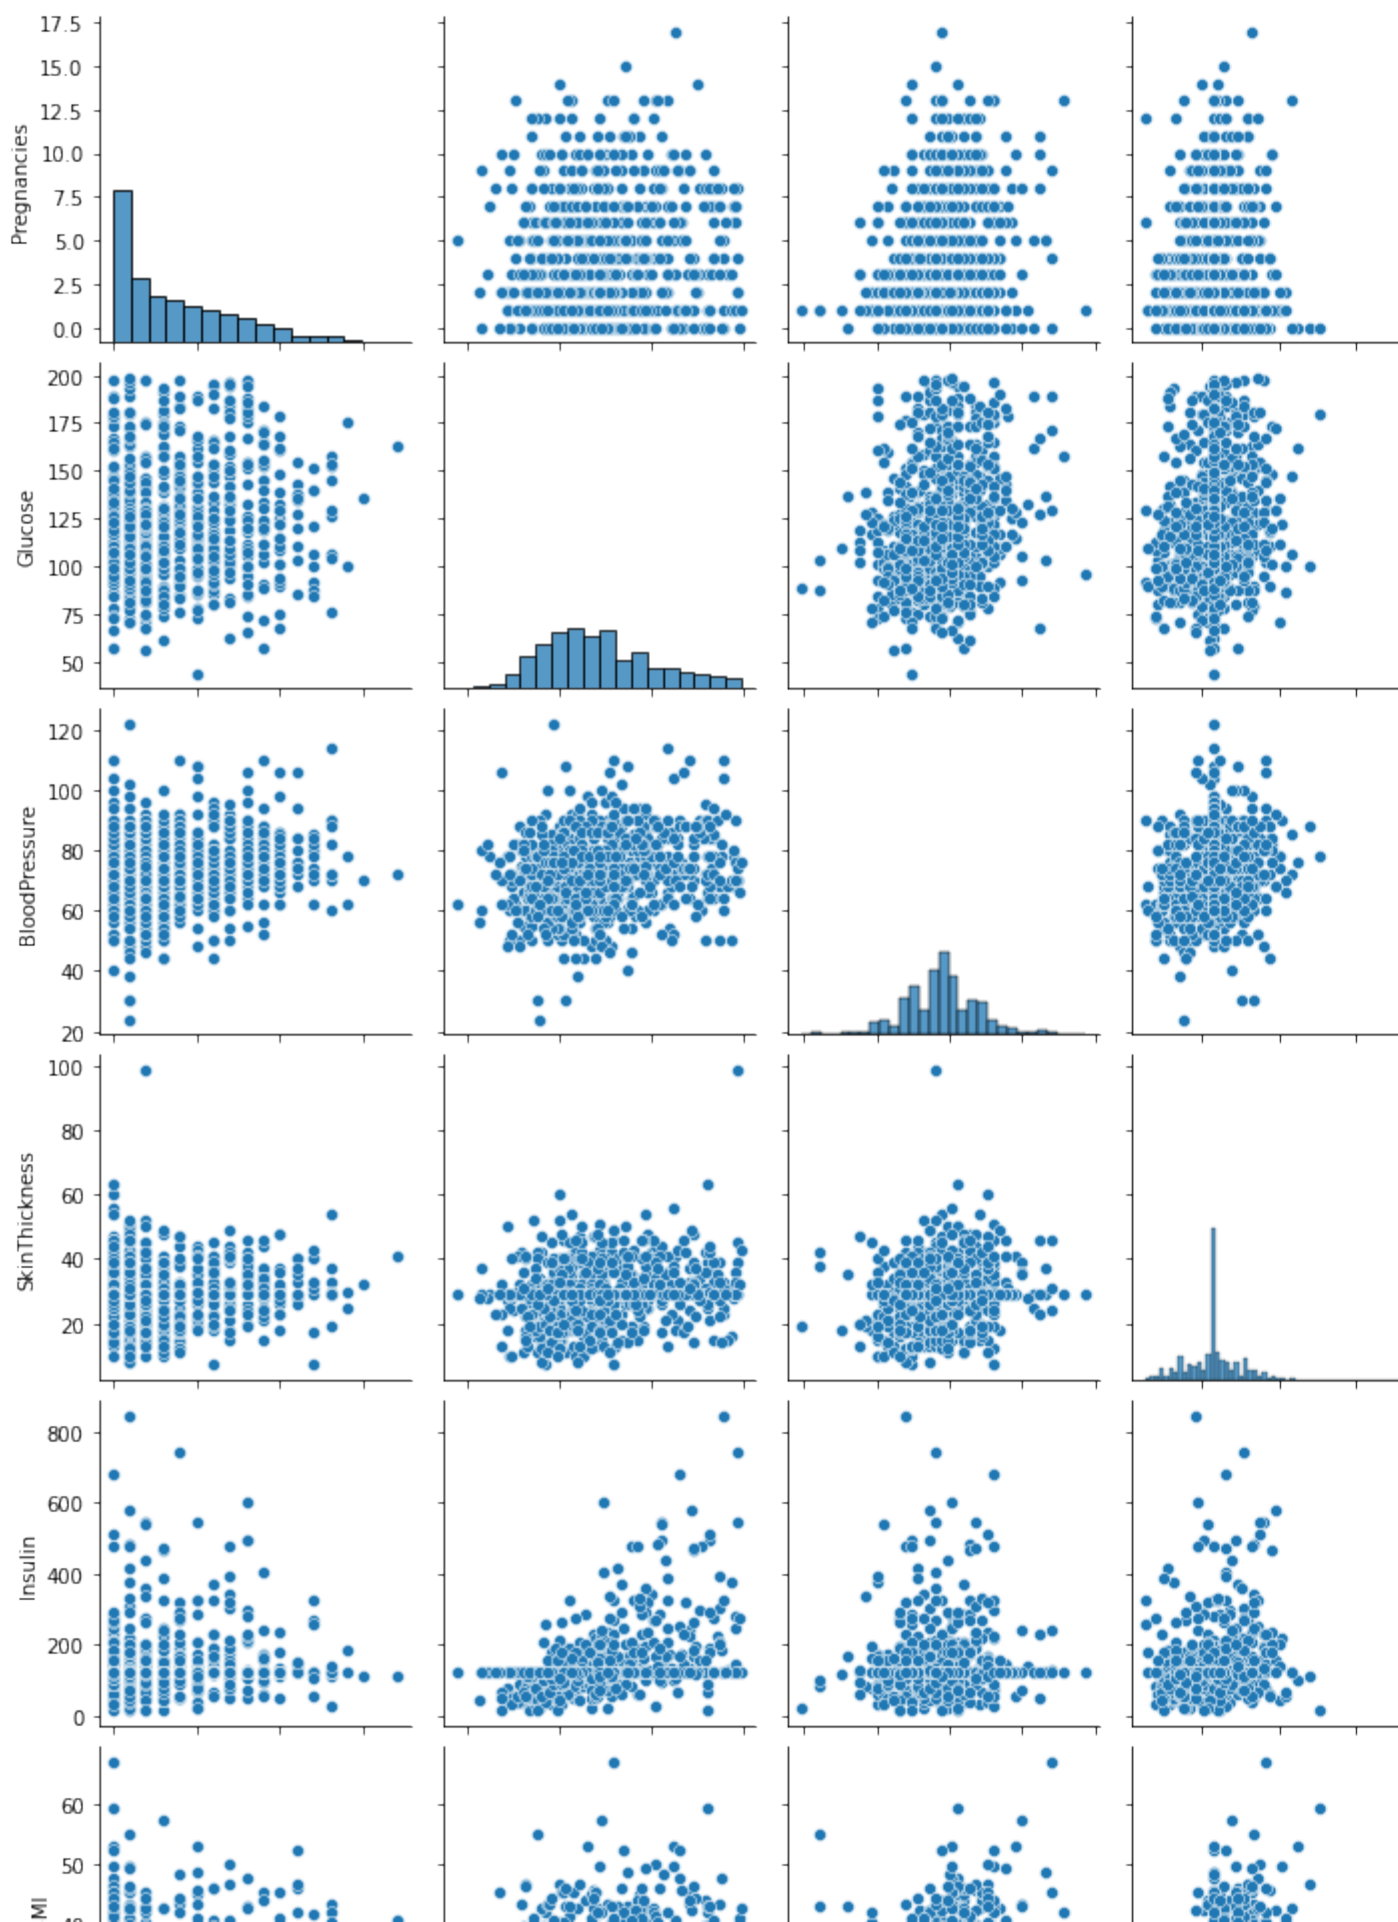
1. Create scatter charts between the pair of variables to understand the relationships. Describe your findings.

In [18]:

```
sns.pairplot(df)
```

Out[18]:

```
<seaborn.axisgrid.PairGrid at 0x7f50206e8210>
```



1. 1. Perform correlation analysis. Visually explore it using a heat map.

In [19]:

```
df.corr()
```

Out[19]:


| | Pregna ncies | Gluc ose | BloodPre ssure | SkinThic kness | Insul in | BMI | DiabetesPedigree Function | Age | Outc ome |
|--------------------------------------|-------------------|--------------|-------------------|-------------------|--------------|--------------|------------------------------|--------------|--------------|
| Pregnancies | 1.00000 0 | 0.128 213 | 0.208615 | 0.081770 | 0.025 047 | 0.021 559 | -0.033523 | 0.544 341 | 0.221 898 |
| Glucose | 0.12821 3 | 1.000 000 | 0.218937 | 0.192615 | 0.419 451 | 0.231 049 | 0.137327 | 0.266 909 | 0.492 782 |
| BloodPressure | 0.20861 5 | 0.218 937 | 1.000000 | 0.191892 | 0.045 363 | 0.281 257 | -0.002378 | 0.324 915 | 0.165 723 |
| SkinThickness | 0.08177 0 | 0.192 615 | 0.191892 | 1.000000 | 0.155 610 | 0.543 205 | 0.102188 | 0.126 107 | 0.214 873 |
| Insulin | 0.02504 7 | 0.419 451 | 0.045363 | 0.155610 | 1.000 000 | 0.180 241 | 0.126503 | 0.097 101 | 0.203 790 |
| BMI | 0.02155 9 | 0.231 049 | 0.281257 | 0.543205 | 0.180 241 | 1.000 000 | 0.153438 | 0.025 597 | 0.312 038 |
| DiabetesPedigree Function | - 0.03352 3 | 0.137 327 | -0.002378 | 0.102188 | 0.126 503 | 0.153 438 | 1.000000 | 0.033 561 | 0.173 844 |
| Age | 0.54434 1 | 0.266 909 | 0.324915 | 0.126107 | 0.097 101 | 0.025 597 | 0.033561 | 1.000 000 | 0.238 356 |
| Outcome | 0.22189 8 | 0.492 782 | 0.165723 | 0.214873 | 0.203 790 | 0.312 038 | 0.173844 | 0.238 356 | 1.000 000 |

In [20]:

```
plt.figure(figsize = (10, 8))  
sns.heatmap(df.corr(),annot = True)
```

Out[20]:

<AxesSubplot:>



Observations show that characteristics like pregnancy, glucose, skinthickness, BMI, and age are more closely associated with outcomes (Glucose as a feature is the most important in this dataset).

Project Task: Week 3

Data Modeling:

1. Devise strategies for model building. It is important to decide the right validation framework. Express your thought process.

In [21]:

```
df.columns
```

Out[21]:

```
Index(['Pregnancies', 'Glucose', 'BloodPressure', 'SkinThickness',  
      'Insulin',
```

```
'BMI', 'DiabetesPedigreeFunction', 'Age', 'Outcome'],
dtype='object')
```

In [22]:

```
features = ['Pregnancies', 'Glucose', 'BloodPressure',
'SkinThickness', 'Insulin', 'BMI', 'DiabetesPedigreeFunction', 'Age']
x= df[features]
x
```

Out[22]:

| | Pregnancies | Glucose | BloodPressure | SkinThickness | Insulin | BMI | DiabetesPedigreeFunction | Age |
|-----|-------------|---------|---------------|---------------|---------|------|--------------------------|-----|
| 0 | 6 | 148 | 72 | 35 | 125 | 33.6 | 0.627 | 50 |
| 1 | 1 | 85 | 66 | 29 | 125 | 26.6 | 0.351 | 31 |
| 2 | 8 | 183 | 64 | 29 | 125 | 23.3 | 0.672 | 32 |
| 3 | 1 | 89 | 66 | 23 | 94 | 28.1 | 0.167 | 21 |
| 4 | 0 | 137 | 40 | 35 | 168 | 43.1 | 2.288 | 33 |
| ... | ... | ... | ... | ... | ... | ... | ... | ... |
| 763 | 10 | 101 | 76 | 48 | 180 | 32.9 | 0.171 | 63 |
| 764 | 2 | 122 | 70 | 27 | 125 | 36.8 | 0.340 | 27 |
| 765 | 5 | 121 | 72 | 23 | 112 | 26.2 | 0.245 | 30 |
| 766 | 1 | 126 | 60 | 29 | 125 | 30.1 | 0.349 | 47 |
| 767 | 1 | 93 | 70 | 31 | 125 | 30.4 | 0.315 | 23 |

768 rows × 8 columns

In [23]:

```
y = df['Outcome'] #target
y
```

Out[23]:

```
0      1
1      0
2      1
3      0
4      1
..
763    0
764    0
765    0
766    1
767    0
Name: Outcome, Length: 768, dtype: int64
```

In [24]:

```
x_train,x_test,y_train,y_test = train_test_split(x,y,test_size=
0.30,random_state=0)
```

In [25]:

```
x_train.shape
```

Out[25]:

```
(537, 8)
```

In [26]:

```
x_test.shape
```

Out[26]:

```
(231, 8)
```

In [27]:

```
y_test
```

Out[27]:

```
661    1
122    0
113    0
14     1
529    0
..
165    1
188    1
334    0
758    0
34     0
```

```
Name: Outcome, Length: 231, dtype: int64
```

1. 1. Apply an appropriate classification algorithm to build a model. Compare various models with the results from KNN algorithm.

Logistic Regression

In [28]:

```
from sklearn.linear_model import LogisticRegression
```

In [29]:

```
logreg = LogisticRegression()
logreg.fit(x_train,y_train)
```

Out[29]:

```
LogisticRegression()
```

In [30]:


```
y_pred= logreg.predict(x_test)
y_pred
```

Out[30]:

```
array([1, 0, 0, 1, 0, 0, 1, 1, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 0, 0, 1,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 0,
1,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
1,
      1, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
1,
      1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
1,
      0, 0, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0,
      0, 1, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      1, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1,
0,
      0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1,
0,
      0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0])
```

In [31]:

```
#Evaluating the performance of the model
from sklearn import metrics
confusion_matrix= metrics.confusion_matrix(y_test,y_pred)
confusion_matrix=
metrics.ConfusionMatrixDisplay(confusion_matrix=confusion_matrix,
display_labels =[False,True])
confusion_matrix.plot()
plt.show()
```



In [32]:

```
#Accuracy
from sklearn.metrics import accuracy_score
```

```
accuracy_score(y_test,y_pred)
```

Out[32]:

```
0.7619047619047619
```

In [33]:

```
#Precision
```

```
from sklearn.metrics import precision_score
precision_score(y_test,y_pred)
```

Out[33]:

```
0.6727272727272727
```

In [34]:

```
#Recall
```

```
from sklearn.metrics import recall_score
recall_score(y_test,y_pred)
```

Out[34]:

```
0.5
```

In [35]:

```
#f1_score
```

```
f1_score=metrics.f1_score(y_test,y_pred)
print(f1_score)
0.5736434108527131
```

Decision Tree

In [36]:

```
from sklearn.tree import DecisionTreeClassifier
dt=DecisionTreeClassifier()
dt.fit(x_train,y_train)
```

Out[36]:

```
DecisionTreeClassifier()
```

In [37]:

```
y_pred1=dt.predict(x_test)
y_pred1
```

Out[37]:

```
array([1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 1,
0,
      0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 0,
1,
      0, 0, 0, 0, 1, 0, 1, 1, 1, 1, 0, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1,
1,
      0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 1, 1, 0, 1, 0,
1,
```



```

1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
0, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
1, 0, 0, 0, 1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0,
0,
0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1,
0,
0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0,
0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 1))


```

In [38]:

```

from sklearn import metrics
confusion_matrix=metrics.confusion_matrix(y_test,y_pred1)
confusion_matrix=metrics.ConfusionMatrixDisplay(confusion_matrix=conf
usion_matrix, display_labels=[False,True])
confusion_matrix.plot()
plt.show()

```



In [39]:

```

#Accuracy
from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred1)

```

Out[39]:

0.7662337662337663

In [40]:

```

#Precision
from sklearn.metrics import precision_score
precision_score(y_test,y_pred1)

```

Out[40]:

0.6219512195121951

In [41]:

```
#Recall
from sklearn.metrics import recall_score
recall_score(y_test,y_pred1)
```

Out[41]:

```
0.6891891891891891
```

In [42]:

```
#f1_score
f1_score=metrics.f1_score(y_test,y_pred1)
print(f1_score)
0.6538461538461539
```

Random Forest

In [43]:

```
from sklearn.ensemble import RandomForestClassifier
rfc=RandomForestClassifier()
rfc.fit(x_train,y_train)
```

Out[43]:

```
RandomForestClassifier()
```

In [44]:


```
y_pred2=rfc.predict(x_test)
y_pred2
```

Out[44]:

```
array([1, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 1, 0, 0, 0,
1,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 1, 1, 0, 1, 1,
0,
      1, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
0,
      0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0,
0,
      0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      1, 0, 1, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 1,
0,
      0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 1, 1, 1, 0, 0, 0, 0, 1])
```

In [45]:

```
from sklearn import metrics
confusion_matrix=metrics.confusion_matrix(y_test,y_pred2)
confusion_matrix=metrics.ConfusionMatrixDisplay(confusion_matrix=confusion_matrix, display_labels=[False,True])
confusion_matrix.plot()
plt.show()
```



In [46]:

```
#Accuracy
from sklearn.metrics import accuracy_score
accuracy_score(y_test,y_pred2)
```

Out[46]:

0.7575757575757576

In [47]:

```
#Precision
from sklearn.metrics import precision_score
precision_score(y_test,y_pred2)
```

Out[47]:

0.6607142857142857

In [48]:

```
#Recall
from sklearn.metrics import recall_score
recall_score(y_test,y_pred2)
```

Out[48]:

0.5

In [49]:

```
#f1_score
```

```
f1_score=metrics.f1_score(y_test,y_pred2)
print(f1_score)
0.5692307692307693
```

KNN

In [50]:

```
from sklearn.neighbors import KNeighborsClassifier
knn=KNeighborsClassifier(n_neighbors=4)
knn.fit(x_train,y_train)
```

Out[50]:

```
KNeighborsClassifier(n_neighbors=4)
```

In [51]:

```
pred=knn.predict(x_test)
pred
```

Out[51]:

```
array([1, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
1,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1,
1,
      1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0,
0,
      1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0,
1,
      0, 0, 1, 0, 1, 1, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 1, 0,
0,
      0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0,
0,
      1, 0, 0, 0, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0,
1,
      0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0,
0,
      0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0, 0])
```


In [52]:

```
from sklearn import metrics

confusion_matrix= metrics.confusion_matrix(y_test,pred)
confusion_matrix=metrics.ConfusionMatrixDisplay(confusion_matrix=conf
usion_matrix,display_labels=[False,True])
confusion_matrix.plot()
```

Out[52]:

<sklearn.metrics._plot.confusion_matrix.ConfusionMatrixDisplay at 0x7f50197877d0>



In [53]:

```
#Accuracy
from sklearn.metrics import accuracy_score
accuracy_score(y_test,pred)
```

Out[53]:

0.7229437229437229

In [54]:

```
#Precision
from sklearn.metrics import precision_score
precision_score(y_test,pred)
```

Out[54]:

0.6086956521739131

In [55]:

```
#Recall
from sklearn.metrics import recall_score
recall_score(y_test,pred)
```

Out[55]:

0.3783783783783784

In [56]:

```
#f1_score
f1_score=metrics.f1_score(y_test,pred)
print(f1_score)
0.4666666666666667
```

Project Task: Week 4

Data Modeling:

1. Create a classification report by analyzing sensitivity, specificity, AUC (ROC curve), etc.
Please be descriptive to explain what values of these parameter you have used.

In [57]:

```
from sklearn.metrics import classification_report, roc_curve,  
roc_auc_score
```

In [58]:

```
print(classification_report(y_test,y_pred)) #Logistic Regression
```

| | precision | recall | f1-score | support | |
|--------------|-----------|--------|----------|---------|-----|
| 0 | | 0.79 | 0.89 | 157 | |
| 1 | | 0.67 | 0.50 | 74 | |
| accuracy | | | | 0.76 | 231 |
| macro avg | | 0.73 | 0.69 | 0.70 | 231 |
| weighted avg | | 0.75 | 0.76 | 0.75 | 231 |

In [59]:

```
print(classification_report(y_test,y_pred1)) #Decision Tree
```

| | precision | recall | f1-score | support | |
|--------------|-----------|--------|----------|---------|-----|
| 0 | | 0.85 | 0.80 | 157 | |
| 1 | | 0.62 | 0.69 | 74 | |
| accuracy | | | | 0.77 | 231 |
| macro avg | | 0.73 | 0.75 | 0.74 | 231 |
| weighted avg | | 0.77 | 0.77 | 0.77 | 231 |

In [60]:

```
print(classification_report(y_test,y_pred2)) #Random Forest
```

| | precision | recall | f1-score | support | |
|--------------|-----------|--------|----------|---------|-----|
| 0 | | 0.79 | 0.88 | 157 | |
| 1 | | 0.66 | 0.50 | 74 | |
| accuracy | | | | 0.76 | 231 |
| macro avg | | 0.72 | 0.69 | 0.70 | 231 |
| weighted avg | | 0.75 | 0.76 | 0.75 | 231 |

In [61]:

```
print(classification_report(y_test,pred)) #KNN
```

| | precision | recall | f1-score | support |
|--|-----------|--------|----------|---------|
|--|-----------|--------|----------|---------|

| | | | | |
|--------------|------|------|------|-----|
| 0 | 0.75 | 0.89 | 0.81 | 157 |
| 1 | 0.61 | 0.38 | 0.47 | 74 |
| accuracy | | | 0.72 | 231 |
| macro avg | 0.68 | 0.63 | 0.64 | 231 |
| weighted avg | 0.71 | 0.72 | 0.70 | 231 |

Therefore Decision Tree is the best model for this prediction since it has an accuracy_score of 0.77.

For Logistic Regression

In [62]:

```
#Using roc_curve() to get the threshold, TPR, and FPR.
fpr,tpr,thresholds =
roc_curve(y_test,logreg.predict_proba(x_test)[:,-1])
print("True Positive Rate - {}, False Positive Rate - {}, Thresholds
- {}".format(tpr,fpr,thresholds))

#For AUC using roc_auc_score()
auc1= roc_auc_score(y_test, logreg.predict(x_test))
print('AUC: %.3f' % auc1)

#Plotting the ROC curve
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
True Positive Rate - [0.          0.01351351 0.04054054 0.04054054
0.05405405 0.05405405
0.12162162 0.12162162 0.14864865 0.14864865 0.18918919 0.18918919
0.28378378 0.28378378 0.32432432 0.32432432 0.41891892 0.41891892
0.43243243 0.43243243 0.45945946 0.45945946 0.47297297 0.47297297
0.48648649 0.48648649 0.54054054 0.54054054 0.55405405 0.55405405
0.58108108 0.58108108 0.59459459 0.59459459 0.62162162 0.62162162
0.66216216 0.66216216 0.68918919 0.68918919 0.71621622 0.71621622
0.74324324 0.74324324 0.75675676 0.75675676 0.78378378 0.78378378
0.81081081 0.81081081 0.83783784 0.83783784 0.85135135 0.85135135
0.86486486 0.86486486 0.87837838 0.87837838 0.89189189 0.89189189
0.90540541 0.90540541 0.91891892 0.91891892 0.94594595 0.94594595
0.95945946 0.95945946 0.97297297 0.97297297 0.98648649 0.98648649
1.          1.          ], False Positive Rate - [0.          0.
0.          0.00636943 0.00636943 0.01273885
0.01273885 0.01910828 0.01910828 0.02547771 0.02547771 0.03184713
0.03184713 0.03821656 0.03821656 0.05095541 0.05095541 0.05732484
0.05732484 0.07006369 0.07006369 0.08280255 0.08280255 0.10828025
0.10828025 0.11464968 0.11464968 0.12738854 0.12738854 0.13375796
0.13375796 0.17197452 0.17197452 0.17834395 0.17834395 0.21019108
0.21019108 0.21656051 0.21656051 0.22929936 0.22929936 0.23566879
0.23566879 0.24840764 0.24840764 0.27388535 0.27388535 0.30573248
0.30573248 0.3566879 0.3566879 0.36942675 0.36942675 0.39490446
```

```


0.39490446 0.42675159 0.42675159 0.43949045 0.43949045 0.44585987
0.44585987 0.47133758 0.47133758 0.5477707 0.5477707 0.66878981
0.66878981 0.68789809 0.68789809 0.7388535 0.7388535 0.82802548
0.82802548 1.          ], Thresholds - [1.97976979 0.97976979
0.96952636 0.96465467 0.95130855 0.93167602
0.87768516 0.86187685 0.84289126 0.8199858 0.79021783 0.77540975
0.71885523 0.71625433 0.70931434 0.70415687 0.62141043 0.62104285
0.61431963 0.59687457 0.57879801 0.55061878 0.54587181 0.52884655
0.5137863 0.51014758 0.49510963 0.4862986 0.47132712 0.46968411
0.45664249 0.43096162 0.41991735 0.41954489 0.41102134 0.39868693
0.39323817 0.39317068 0.39003323 0.38116884 0.37154483 0.36954273
0.35394002 0.34148568 0.34102858 0.3267451 0.31304148 0.29234675
0.29182313 0.26468055 0.26365187 0.25865693 0.25526253 0.25129361
0.24875032 0.24288912 0.2417289 0.23000866 0.22963736 0.22266081
0.2222047 0.21113181 0.20706676 0.17565968 0.16860768 0.13140044
0.13034169 0.12643356 0.12560678 0.11482271 0.11389767 0.0803826
0.07768083 0.03802583]

```

AUC: 0.693

Out[62]:

Text(0, 0.5, 'True Positive Rate')



For Decision Tree

In [63]:

```

#Using roc_curve() to get the threshold, TPR, and FPR.
fpr,tpr,thresholds = roc_curve(y_test,dt.predict_proba(x_test)[: ,1])
print("True Positive Rate - {}, False Positive Rate - {}, Thresholds
- {}".format(tpr,fpr,thresholds))

#For AUC using roc_auc_score()
auc2= roc_auc_score(y_test, dt.predict(x_test))
print('AUC: %.3f' % auc2)

#Ploting the ROC curve


```



```
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
True Positive Rate - [0.          0.68918919 1.          ], False
Positive Rate - [0.          0.19745223 1.          ], Thresholds - [2.
1. 0.]
AUC: 0.746
```

Out[63]:

```
Text(0, 0.5, 'True Positive Rate')
```



For Random Forest

In [64]:

```
#Using roc_curve() to get the threshold, TPR, and FPR.
fpr,tpr,thresholds = roc_curve(y_test,rfc.predict_proba(x_test)[:,-1])
print("True Positive Rate - {}, False Positive Rate - {}, Thresholds
- {}".format(tpr,fpr,thresholds))

#For AUC using roc_auc_score()
auc3= roc_auc_score(y_test, rfc.predict(x_test))
print('AUC: %.3f' % auc3)

#Plotting the ROC curve
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
True Positive Rate - [0.          0.01351351 0.04054054 0.06756757
0.08108108 0.10810811
0.12162162 0.14864865 0.14864865 0.17567568 0.2027027  0.22972973
0.24324324 0.25675676 0.28378378 0.28378378 0.2972973  0.2972973
0.33783784 0.35135135 0.35135135 0.36486486 0.36486486 0.39189189
0.41891892 0.43243243 0.45945946 0.45945946 0.54054054 0.58108108
```


```

0.60810811 0.62162162 0.63513514 0.63513514 0.63513514 0.66216216
0.67567568 0.72972973 0.75675676 0.75675676 0.75675676 0.75675676
0.78378378 0.81081081 0.82432432 0.82432432 0.83783784 0.86486486
0.86486486 0.86486486 0.89189189 0.90540541 0.90540541 0.90540541
0.90540541 0.91891892 0.94594595 0.97297297 0.97297297 0.97297297
0.97297297 0.98648649 0.98648649 0.98648649 1. 1.
1. 1. 1. 1. 1. 1. ],
False Positive Rate - [0. 0. 0. 0. 0. 0.
0.
0. 0. 0.01273885 0.01910828 0.01910828 0.01910828
0.03184713 0.03821656 0.03821656 0.04458599 0.04458599 0.05732484
0.05732484 0.06369427 0.07643312 0.08280255 0.08917197 0.08917197
0.10191083 0.10191083 0.10191083 0.12101911 0.12101911 0.12738854
0.14012739 0.14649682 0.15923567 0.1656051 0.20382166 0.21019108
0.21019108 0.22292994 0.22292994 0.22929936 0.24203822 0.24840764
0.27388535 0.28025478 0.29936306 0.31847134 0.32484076 0.36305732
0.38216561 0.38853503 0.39490446 0.42675159 0.43949045 0.47133758
0.49044586 0.50955414 0.53503185 0.55414013 0.58598726 0.59872611
0.63057325 0.66878981 0.69426752 0.70700637 0.74522293 0.78343949
0.8343949 0.84713376 0.9044586 0.91082803 0.94904459 1. ],
Thresholds - [1.96 0.96 0.89 0.88 0.86 0.83 0.82 0.81 0.8 0.79 0.78
0.76 0.75 0.74
0.73 0.72 0.69 0.68 0.67 0.66 0.65 0.64 0.63 0.62 0.57 0.55 0.54
0.53
0.5 0.49 0.48 0.47 0.46 0.45 0.43 0.42 0.41 0.39 0.38 0.36 0.35
0.34
0.32 0.31 0.29 0.28 0.27 0.25 0.24 0.23 0.22 0.21 0.19 0.18 0.17
0.16
0.15 0.14 0.13 0.12 0.11 0.1 0.09 0.08 0.07 0.06 0.05 0.04 0.03
0.02
0.01 0. ]
AUC: 0.689

```

Out[64]:

```
Text(0, 0.5, 'True Positive Rate')
```



For KNN

In [65]:


```
#Using roc_curve() to get the threshold, TPR, and FPR.
fpr,tpr,thresholds = roc_curve(y_test,knn.predict_proba(x_test)[:,-1])
print("True Positive Rate - {}, False Positive Rate - {}, Thresholds
- {}".format(tpr,fpr,thresholds))

#For AUC using roc_auc_score()
auc= roc_auc_score(y_test, knn.predict(x_test))
print('AUC: {:.3f}' % auc)

#Plotting the ROC curve
plt.plot([0, 1], [0, 1], linestyle='--')
plt.plot(fpr, tpr, marker='.')
plt.xlabel("False Positive Rate")
plt.ylabel("True Positive Rate")
True Positive Rate - [0.012162162 0.37837838 0.58108108
0.85135135 1.], False Positive Rate - [0.001910828
0.11464968 0.28025478 0.55414013 1.], Thresholds - [2. 1.
0.75 0.5 0.25 0. ]
AUC: 0.632
```

Out[65]:

Text(0, 0.5, 'True Positive Rate')




Create a dashboard


HealthCare

Correlation Analysis


| | Age Rank | | | | | | | | | | | | |
|---------------------|----------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|-------|--|
| | 20-25 | 26-30 | 31-35 | 36-40 | 41-45 | 46-50 | 51-55 | 56-60 | 61-65 | 66-70 | 71-75 | 81-85 | |
| Avg. BMI | 30.4 | 33.0 | 32.8 | 33.0 | 35.3 | 32.9 | 31.8 | 30.2 | 29.9 | 27.5 | 19.6 | 25.9 | |
| Avg. Blood Pressure | 63.8 | 68.0 | 68.8 | 71.8 | 73.3 | 77.9 | 81.9 | 77.5 | 76.0 | 80.7 | 0.0 | 74.0 | |
| Avg. Glucose | 110.7 | 120.3 | 124.2 | 128.3 | 125.1 | 124.5 | 143.2 | 138.3 | 136.4 | 139.0 | 119.0 | 134.0 | |
| Avg. Insulin | 84.3 | 84.3 | 92.8 | 59.5 | 56.7 | 67.6 | 109.9 | 149.5 | 26.4 | 0.0 | 0.0 | 60.0 | |
| Avg. Skin Thickness | 22.0 | 21.3 | 20.1 | 21.0 | 18.9 | 20.4 | 16.3 | 18.7 | 20.0 | 1.6 | 0.0 | 33.0 | |



Bubble Chart for BP



Age wise Diabetic/Non-Diabetic



Histogram for BP

